

Machine Learning with R and H2O

SPENCER AIELLO ERIC ECKSTRAND ANQI FU
MARK LANDRY JESSICA LANFORD PATRICK ABOYOUN

<http://h2o.ai/resources/>

October 2015: Third Edition

Machine Learning with R and H2O
by Spencer Aiello,
Eric Eckstrand, Anqi Fu, Mark Landry,
Jessica Lanford & Patrick Aboyoun

Published by H2O.ai, Inc.
2307 Leghorn St.
Mountain View, CA 94043

©2015 H2O.ai, Inc. All Rights Reserved.

September 2015: Third Edition

Photos by ©H2O.ai, Inc.

All copyrights belong to their respective owners.
While every precaution has been taken in the
preparation of this book, the publisher and
authors assume no responsibility for errors or
omissions, or for damages resulting from the
use of the information contained herein.

Printed in the United States of America.

Contents

1	Introduction	5
2	What is H2O?	6
3	Installation	7
3.1	Installing R	7
3.2	Installing H2O from R	8
3.3	Citation	8
4	H2O Initialization	9
4.1	Launching from R	9
4.2	Launching from the Command Line	10
4.3	Launching on Hadoop	11
4.4	Launching on an EC2	11
4.5	Checking Cluster Status	11
5	Data Preparation in R	12
5.1	Notes	12
5.2	Demo: Creating Aggregates from Split Data	13
6	Models	14
6.1	Supervised Learning	14
6.2	Unsupervised Learning	15
6.3	Modeling Constructs	15
7	Demo: GLM	15
8	Data Manipulation in R	17
8.1	Importing Files	17
8.2	Uploading Files	18
8.3	Finding Factors	18
8.4	Converting to Factors	18
8.5	Converting Data Frames	19
8.6	Transferring Data Frames	19
8.7	Renaming Data Frames	20
8.8	Viewing Column Names	20
8.9	Getting Minimum and Maximum Values	21
8.10	Getting Quantiles	21
8.11	Summarizing Data	22
8.12	Summarizing Data in a Table	22
8.13	Generating Random Numbers	22
8.14	Splitting Frames	23

8.15	Getting Frames	24
8.16	Getting Models	24
8.17	Listing H2O Objects	25
8.18	Removing H2O Objects	25
8.19	Adding Functions	25
9	Running Models	25
9.1	Gradient Boosted Models (GBM)	26
9.2	Generalized Linear Models (GLM)	26
9.3	K-Means	27
9.4	Principal Components Analysis (PCA)	28
9.5	Predictions	28
10	Appendix: Commands	30
10.1	Data Set Operations	30
10.2	General Data Operations	31
10.3	Methods from Group Generics	32
10.4	Other Aggregations	34
10.5	Data Munging	35
10.6	Data Modeling	36
10.7	H2O Cluster Operations	37
11	Authors	40
12	References	41

1 Introduction

This documentation describes how to use H2O in the R environment. More information on H2O's system and algorithms (as well as R user documentation) is available at the H2O website at <http://docs.h2o.ai>.

R uses a REST API to connect to H2O, so you can specify the IP address and port number of the H2O instance in the R environment to use H2O in R or launch H2O from R. Data sets are not directly transmitted through the REST API. Instead, sending a command (for example, an HDFS path to import a data set) either through the browser or the REST API performs the specified task.

The data set is then assigned an identifier, which uses the .hex file type in H2O, that is used as a reference in commands to the web server. After preparing the dataset for modeling by defining significant data and removing insignificant data, H2O creates a model representing the results of the data analysis. These models are assigned IDs that are used as references in commands. One of the most popular models for data analysis is GLM.

GLM estimates regression models for outcomes following exponential distributions in general. In addition to the Gaussian (i.e. normal) distribution, these include binomial, gamma, Poisson, and Tweedie distributions. Each serves a different purpose, and depending on distribution and link function, can be used for prediction or classification.

This booklet demonstrates the use of H2O's implementation of GLM in an R environment. For more information on GLM, there is an additional booklet (similar to this one) available at <http://h2o.ai/resources/>.

H2O supports Spark, YARN, and all versions of Hadoop. Hadoop is a scalable open-source file system that uses clusters for distributed storage and dataset processing. Depending on the size of your data, H2O can run on your desktop or scale using multiple nodes with Hadoop, an EC2 cluster, or S3 storage.

H2O nodes run as JVM invocations on Hadoop nodes. For performance reasons, we recommend that you do not run an H2O node on the same hardware as the Hadoop NameNode.

Because H2O nodes run as mapper tasks in Hadoop, administrators can view them in the normal JobTracker and TaskTracker frameworks, providing process-level (i.e. JVM instance-level) visibility.

H2O helps R users make the leap from laptop-based processing to large-scale environments. Hadoop lets H2O users scale their data processing capabilities based on their current needs. Using H2O, R, and Hadoop, you can create a complete end-to-end data analysis solution.

This document describes the four steps of data analysis with H2O:

1. installing H2O
2. preparing your data for modeling (data munging)
3. creating a model using simple but powerful machine learning algorithms
4. scoring your models

2 What is H2O?

H2O is fast, scalable, open-source machine learning and deep learning for smarter applications. With H2O, enterprises like PayPal, Nielsen Catalina, Cisco, and others can use all their data without sampling to get accurate predictions faster. Advanced algorithms such as deep learning, boosting, and bagging ensembles are built-in to help application designers create smarter applications through elegant APIs. Some of our initial customers have built powerful domain-specific predictive engines for recommendations, customer churn, propensity to buy, dynamic pricing, and fraud detection for the insurance, healthcare, telecommunications, ad tech, retail, and payment systems industries.

Using in-memory compression, H2O handles billions of data rows in-memory, even with a small cluster. To make it easier for non-engineers to create complete analytic workflows, H2O's platform includes interfaces for R, Python, Scala, Java, JSON, and CoffeeScript/JavaScript, as well as a built-in web interface, Flow. H2O is designed to run in standalone mode, on Hadoop, or within a Spark Cluster, and typically deploys within minutes.

H2O includes many common machine learning algorithms, such as generalized linear modeling (linear regression, logistic regression, etc.), Naïve Bayes, principal components analysis, k-means clustering, and others. H2O also implements best-in-class algorithms at scale, such as distributed random forest, gradient boosting, and deep learning. Customers can build thousands of models and compare the results to get the best predictions.

H2O is nurturing a grassroots movement of physicists, mathematicians, and computer scientists to herald the new wave of discovery with data science by collaborating closely with academic researchers and industrial data scientists. Stanford university giants Stephen Boyd, Trevor Hastie, Rob Tibshirani advise the H2O team on building scalable machine learning algorithms. With hundreds of meetups over the past three years, H2O has become a word-of-mouth phenomenon, growing amongst the data community by a hundred-fold, and is now used by 30,000+ users and is deployed using R, Python, Hadoop, and Spark in 2000+ corporations.

Try it out

- Download H2O directly at <http://h2o.ai/download>.
- Install H2O's R package from CRAN at <https://cran.r-project.org/web/packages/h2o/>.
- Install the Python package from PyPI at <https://pypi.python.org/pypi/h2o/>.

Join the community

- To learn about our meetups, training sessions, hackathons, and product updates, visit <http://h2o.ai>.
- Visit the open source community forum at <https://groups.google.com/d/forum/h2ostream>.
- Join the chat at <https://gitter.im/h2oai/h2o-3>.

3 Installation

H2O requires Java; if you do not already have Java installed, install it from <https://java.com/en/download/> before installing H2O.

To use H2O with R, start H2O outside of R and connect to it, or launch H2O from R. However, if you launch H2O from R and close the R session, the H2O session closes as well. The H2O session directs R to the datasets and models located in H2O.

The following sections describe:

- installing R
- installing H2O from R

3.1 Installing R

To download R:

1. Go to <http://cran.r-project.org/mirrors.html>.
2. Select your closest local mirror.
3. Select your operating system (Linux, OS X, or Windows).
4. Depending on your OS, download the appropriate file, along with any required packages.

5. When the download is complete, unzip the file and install.

3.2 Installing H2O from R

Example in R

To load a recent H2O package from CRAN, run:

```
1 install.packages("h2o")
```

Note: The version of H2O in CRAN is often one release behind the current version.

For the latest recommended version, download the latest stable H2O-3 build from the H2O download page:

1. Go to <http://h2o.ai/download>.
2. Choose the latest stable H2O-3 build.
3. Click the “Install in R” tab.
4. Copy and paste the commands into your R session.

Example in R

After H2O is installed on your system, verify the installation completed successfully by initializing H2O:

```
1 library(h2o)
2
3 #Start H2O on your local machine using all available
   cores
4 #(By default, CRAN policies limit use to only 2 cores)
5 h2o.init(nthreads = -1)
6
7 #Get help
8 ?h2o.glm
9 ?h2o.gbm
10
11 #Show a demo
12 # demo(h2o.glm)
13 # demo(h2o.gbm)
```

3.3 Citation

To cite this booklet, use the following:

Aiello, S., Eckstrand, E., Fu, A., Landry, M., Lanford, J., and Aboyoun, P. (Aug. 2015) *Fast Scalable R with H2O*. <http://h2o.ai/resources/>.

4 H2O Initialization

This section describes how to launch H2O:

- from R
- from the command line
- on Hadoop
- on an EC2 cluster

4.1 Launching from R

To specify the number of CPUs for the H2O session, use the `nthreads` = parameter in the `h2o.init` command. `-2` uses the CRAN default of 2 CPUs. `-1` uses all CPUs on the host, which is strongly recommended. To use a specific number of CPUs, enter a positive integer.

To specify the maximum amount of memory for the H2O session, use the `max_mem_size` parameter in the `h2o.init` command. The value must be a multiple of 1024 greater than 2MB. Append the letter `m` or `M` to indicate megabytes, or `g` or `G` to indicate gigabytes.

If you do not specify a value for `max_mem_size` when you run `h2o.init`, the default heap size of the H2O instance running on 32-bit Java is 1g.

For best performance, the allocated memory should be 4x the size of your data, but never more than the total amount of memory on your computer. For larger data sets, we recommend running on a server or service with more memory available for computing.

H2O checks the Java version and suggests an upgrade if you are running 32-bit Java. On 64-bit Java, the heap size is 1/4 of the total memory available on the machine.

Example in R

To launch H2O locally from R, run the following in R:

```
1 library(h2o)
2 # Starts H2O using localhost IP, port 54321, all CPUs,
  # and 4g of memory
3 h2o.init(ip = 'localhost', port = 54321, nthreads= -1,
  max_mem_size = '4g')
```

After successfully launching, R displays output similar to the following example:

```
1 Successfully connected to http://localhost:54321
2 R is connected to H2O cluster:
3   H2O cluster uptime:      11 minutes 35 seconds
4   H2O cluster version:    2.7.0.1497
5   H2O cluster name:       H2O_started_from_R
6   H2O cluster total nodes: 1
7   H2O cluster total memory: 3.56 GB
8   H2O cluster total cores: 8
9   H2O cluster allowed cores: 8
10  H2O cluster healthy:    TRUE
```

Example in R

To launch H2O locally with default initialization arguments, use the following:

```
1 h2o.init()
```

Example in R

To connect to an established H2O cluster (in a multi-node Hadoop environment, for example) specify the IP address and port number for the established cluster using the `ip` and `port` parameters in the `h2o.init()` command.

```
1 h2o.init(ip = "123.45.67.89", port = 54321)
```

4.2 Launching from the Command Line

A simple way to launch H2O from the command line is to download the H2O zip file from the H2O download page. Unzip and launch H2O with the following:

```
1 unzip h2o-3.2.0.1-*.zip
2 cd h2o-3.2.0.1-*
3 java -jar h2o.jar
```

See the H2O Documentation for additional JVM and H2O command line options. After launching the H2O instance, connect to it from R with `h2o.init()` as described above.

4.3 Launching on Hadoop

To launch H2O nodes and form a cluster on the Hadoop cluster, run:

```
1 hadoop jar h2odriver.jar -nodes 1 -mapperXmx 6g -
  output hdfsOutputDirName
```

- You must launch the Hadoop-specific H2O driver jar (`h2odriver.jar`) for your Hadoop distribution. Specific driver jar files are available for the following Hadoop versions:

- cdh5.2	- hdp2.1	- mapr3.1.1
- cdh5.3	- hdp2.2	- mapr4.0.1
- cdh5.4.2		- mapr5.0

- The above command launches exactly one 6g node of H2O; however, we recommend launching the cluster with 4 times the memory of your data file.
- `mapperXmx` is the mapper size or the amount of memory allocated to each node.
- `nodes` is the number of nodes requested to form the cluster.
- `output` is the name of the directory created each time a H2O cloud is created so it is necessary for the name to be unique each time it is launched.

4.4 Launching on an EC2

Note: If you would like to try out H2O on an EC2 cluster, <http://play.h2o.ai> is the easiest way to get started. H2O Play provides access to a temporary cluster managed by H2O.

4.5 Checking Cluster Status

Example in R

To check the status and health of the H2O cluster, use `h2o.clusterInfo()`.

```
1 library(h2o)
2 h2o.init()
3 h2o.clusterInfo()
```

An easy-to-read summary of information about the cluster displays.

```
1
2 R is connected to H2O cluster:
3   H2O cluster uptime:      43 minutes 43 seconds
4   H2O cluster version:    2.7.0.1497
5   H2O cluster name:       H2O_started_from_R
6   H2O cluster total nodes: 1
7   H2O cluster total memory: 3.56 GB
8   H2O cluster total cores: 8
9   H2O cluster allowed cores: 8
10  H2O cluster healthy:     TRUE
```

5 Data Preparation in R

The following section contains information about data preparation (also known as data munging) and some of the tools and methods available in H2O, as well as a data training example.

5.1 Notes

- Although it may seem like you are manipulating the data in R, once the data has been passed to H2O, all data munging occurs in the H2O instance. The information is passed to R through JSON APIs, so some functions may not have another method.
- You are limited by the total amount of memory allocated to the H2O instance, not by R's ability to handle data. To process large data sets, make sure to allocate enough memory. For more information, refer to **Launching from R**.
- You can manipulate datasets with thousands of factor levels using H2O in R, so if you ask H2O to display a table in R with information from high cardinality factors, the results may overwhelm R's capacity.
- To manipulate data in R and not in H2O, use `as.data.frame()`, `as.h2o()`, and `str()`.

- `as.data.frame()` converts an H2O data frame into an R data frame. If your request exceeds the amount of data supported by R, the R session will crash. If possible, we recommend only taking subsets of the entire data set (the necessary data columns or rows) instead of the whole data set.
- `as.h2o()` transfers data from R to the H2O instance. For successful data transfer, we recommend confirming enough memory is allocated to the H2O instance.
- `str.H2OFrame()` returns the elements of the new object to confirm that the data transferred correctly. It's a good way to verify there were no data loss or conversion issues.

5.2 Demo: Creating Aggregates from Split Data

The following section depicts an example of creating aggregates for data training using `ddply()`. Using this method, you can split your dataset and apply a function to the subsets.

Example in R

To apply a user-specified function to each subset of an H2O dataset and combine the results, use `ddply()`, with the name of the H2O object, the variable name, and the function in the parentheses.

```

1 library(h2o)
2 h2o.init(nthreads = -1)
3
4 # Import iris dataset to H2O
5 irisPath = system.file("extdata", "iris_wheader.csv",
6   package = "h2o")
7
8 # Apply function to groups by class of flower
9 # uses h2o's ddply, since iris.hex is an H2OFrame
10 object
11 res = h2o.ddply(iris.hex, "class", function(df) { sum(
12   df[,1], na.rm = T)/nrow(df) })
13 head(res)

```

6 Models

The following section describes the features and functions of some common models available in H2O. For more information about running these models in R using H2O, refer to **Running Models**.

H2O supports the following models:

- Deep Learning
- Naïve Bayes
- Principal Components Analysis (PCA)
- K-means
- Generalized Linear Models (GLM)
- Gradient Boosted Regression (GBM)
- Distributed Random Forest (DRF)

The list is growing quickly, so check www.h2o.ai to see the latest additions. The following list describes some common model types and features.

6.1 Supervised Learning

Generalized Linear Models (GLM): Provides flexible generalization of ordinary linear regression for response variables with error distribution models other than a Gaussian (normal) distribution. GLM unifies various other statistical models, including Poisson, linear, logistic, and others when using ℓ_1 and ℓ_2 regularization.

Distributed Random Forest: Averages multiple decision trees, each created on different random samples of rows and columns. It is easy to use, non-linear, and provides feedback on the importance of each predictor in the model, making it one of the most robust algorithms for noisy data.

Gradient Boosting (GBM): Produces a prediction model in the form of an ensemble of weak prediction models. It builds the model in a stage-wise fashion and is generalized by allowing an arbitrary differentiable loss function. It is one of the most powerful methods available today.

Deep Learning: Models high-level abstractions in data by using non-linear transformations in a layer-by-layer method. Deep learning is an example of supervised learning, which can use unlabeled data that other algorithms cannot.

Naïve Bayes: Generates a probabilistic classifier that assumes the value of a particular feature is unrelated to the presence or absence of any other feature, given the class variable. It is often used in text categorization.

6.2 Unsupervised Learning

K-Means: Reveals groups or clusters of data points for segmentation. It clusters observations into k -number of points with the nearest mean.

Anomaly Detection: Identifies the outliers in your data by invoking the deep learning autoencoder, a powerful pattern recognition model.

6.3 Modeling Constructs

Grid Search: Performs standard hyper-parameter optimization to simplify model configuration.

After creating a model, use it to make predictions. For more information about predictions, refer to **Predictions**.

7 Demo: GLM

The following demo demonstrates how to:

1. Import a file
2. Define significant data
3. View data
4. Create testing and training sets using sampling
5. Define the model
6. Display the results

Example in R

```
1 # Import dataset and display summary
2 library(h2o)
3 h2o.init()
4 airlinesURL = "https://s3.amazonaws.com/h2o-airlines-
  unpacked/allyears2k.csv"
5 airlines.hex = h2o.importFile(path = airlinesURL,
  destination_frame = "airlines.hex")
6 summary(airlines.hex)
7
8 # View quantiles and histograms
9 #high_na_columns = h2o.ignoreColumns(data = airlines.
  hex)
10 quantile(x = airlines.hex$ArrDelay, na.rm = TRUE)
11 h2o.hist(airlines.hex$ArrDelay)
```

```
12
13 # Find number of flights by airport
14 originFlights = h2o.ddply(airlines.hex, 'Origin', now
   )
15 originFlights.R = as.data.frame(originFlights)
16
17 # Find number of cancellations per month
18 flightsByMonth = h2o.ddply(airlines.hex,"Month", now)
19 flightsByMonth.R = as.data.frame(originFlights)
20
21 # Find months with the highest cancellation ratio
22 which(colnames(airlines.hex)== "Cancelled")
23 fun = function(df) {sum(df[,22])}
24 cancellationsByMonth = h2o.ddply(airlines.hex,"Month",
   fun)
25 cancellation_rate = cancellationsByMonth$C1/
   flightsByMonth$C1
26 rates_table = cbind(flightsByMonth$Month, cancellation
   _rate)
27 rates_table.R = as.data.frame(rates_table)
28
29 # Construct test and train sets using sampling
30 airlines.split = h2o.splitFrame(data = airlines.hex,
   ratios = 0.85)
31 airlines.train = airlines.split[[1]]
32 airlines.test = airlines.split[[2]]
33
34 # Display a summary using table-like functions
35 h2o.table(airlines.train$Cancelled)
36 h2o.table(airlines.test$Cancelled)
37
38 # Set predictor and response variables
39 Y = "IsDepDelayed"
40 X = c("Origin", "Dest", "DayofMonth", "Year", "
   UniqueCarrier", "DayOfWeek", "Month", "DepTime", "
   ArrTime", "Distance")
41 # Define the data for the model and display the
   results
42 airlines.glm <- h2o.glm(training_frame=airlines.train,
   x=X, y=Y, family = "binomial", alpha = 0.5)
43 # View model information: training statistics,
   performance, important variables
44 summary(airlines.glm)
```



```

45 |
46 | # Predict using GLM model
47 | pred = h2o.predict(object = airlines.glm, newdata =
    |   airlines.test)
48 | # Look at summary of predictions: probability of TRUE
    | class (p1)
49 | summary(pred$p1)

```

8 Data Manipulation in R

The following section describes some common R commands. For a complete command list, including parameters, refer to http://h2o-release.s3.amazonaws.com/h2o/latest_stable_Rdoc.html. For additional help within R's Help tab, precede the command with a question mark (for example, ?h2o) for suggested commands containing the search terms. For more information on a command, precede the command with two question marks (??h2o).

8.1 Importing Files

Example in R

The H2O package consolidates all of the various supported import functions using `h2o.importFile()`. There are a few ways to import files shown in the following examples:

```

1 | #To import small iris data file from H2O's package:
2 | irisPath = system.file("extdata", "iris.csv", package=
    |   "h2o")
3 | iris.hex = h2o.importFile(path = irisPath, destination
    |   _frame = "iris.hex")
4 |
5 | #To import an entire folder of files as one data
    | object:
6 | # pathToFolder = "/Users/data/airlines/"
7 | # airlines.hex = h2o.importFile(path = pathToFolder,
    | destination_frame = "airlines.hex")
8 |
9 | #To import from HDFS and connect to H2O in R using the
    | IP and port of an H2O instance running on your
    | Hadoop cluster:
10 | # h2o.init(ip= <IPAddress>, port =54321, nthreads =
    | -1)

```

```
11 # pathToData = "hdfs://mr-0xd6.h2oai.loc/datasets/  
    airlines_all.csv"  
12 # airlines.hex = h2o.importFile(path = pathToData,  
    destination_frame = "airlines.hex")
```

8.2 Uploading Files

Example in R

To upload a file in a directory local to your H2O instance, use `h2o.importFile()`. `h2o.uploadFile()` uploads data local to your H2O instance as well as uploading data local to your R session. In the parentheses, specify the H2O reference object in R and the complete URL or normalized file path for the file.

```
1 irisPath = system.file("extdata", "iris.csv", package=  
    "h2o")  
2 iris.hex = h2o.uploadFile(path = irisPath, destination  
    _frame = "iris.hex")
```

8.3 Finding Factors

Example in R

To determine if any column contains categorical data (also known as a factor), use `h2o.anyFactor()`, with the R reference object in the parentheses.

```
1 irisPath = system.file("extdata", "iris_wheader.csv",  
    package="h2o")  
2 iris.hex = h2o.importFile(path = irisPath)  
3 h2o.anyFactor(iris.hex)
```

8.4 Converting to Factors

Example in R

To convert an integer into a non-ordered factor (also called an enum or categorical), use `as.factor()` with the name of the R reference object in parentheses, followed by the number of the column to convert in brackets.

```
1 # Import prostate data  
2 prosPath <- system.file("extdata", "prostate.csv",  
    package="h2o")  
3 prostate.hex <- h2o.importFile(path = prosPath)
```

```
4
5 # Converts column 4 (RACE) to an enum
6 is.factor(prostate.hex[,4])
7
8 prostate.hex[,4] <- as.factor(prostate.hex[,4])
9 is.factor(prostate.hex[,4])
10
11 # Summary will return a count of the factors
12 summary(prostate.hex[,4])
```

8.5 Converting Data Frames

Example in R

To convert an H2O parsed data object into an R data frame that can be manipulated using R commands, use `as.data.frame()` with the name of the R reference object in the parentheses. **Caution:** While this can be very useful, be careful when using this command to convert H2O parsed data objects. H2O can easily handle data sets that are often too large to be handled equivalently well in R.

```
1 # Creates object that defines path
2 prosPath <- system.file("extdata", "prostate.csv",
3   package="h2o")
4 # Imports data set
5 prostate.hex = h2o.importFile(path = prosPath,
6   destination_frame="prostate.hex")
7
8 # Converts current data frame (prostate data set) to
9   an R data frame
10 prostate.R <- as.data.frame(prostate.hex)
11 # Displays a summary of data frame where the summary
12   was executed in R
13 summary(prostate.R)
```

8.6 Transferring Data Frames

Example in R

To transfer a data frame from the R environment to the H2O instance, use `as.h2o()`. In the parentheses, specify the object in the R environment to convert to an H2O object. Optionally, include the name of the destination frame in H2O. Precede the destination frame name with

`destination_frame =` and enclose the name in quotes as in the following example.

```
1 # Import the iris data into H2O
2 data(iris)
3 iris
4
5 # Converts R object "iris" into H2O object "iris.hex"
6 iris.hex = as.h2o(iris, destination_frame= "iris.hex")
7
8 head(iris.hex)
```

8.7 Renaming Data Frames

Example in R

To rename a dataframe on the server running H2O for a data set manipulated in R, use `h2o.assign()`. In the following example, the prostate data set was uploaded to the H2O instance and the data was manipulated to remove outliers. `h2o.assign()` saves the new data set on the H2O server so it can be analyzed using H2O without overwriting the original data set.

```
1 prosPath <- system.file("extdata", "prostate.csv",
2   package="h2o")
3 prostate.hex<-h2o.importFile(path = prosPath)
4
5 ## Assign a new name to prostate dataset in the KV
6   store
7 h2o.ls()
8
9 prostate.hex <- h2o.assign(data = prostate.hex, key =
10   "myNewName")
11 h2o.ls()
```

8.8 Viewing Column Names

Example in R

To view a list of the column names in the data set, use `colnames()` or `names()` with the name of the R reference object in the parentheses.

```
1 ##Displays the titles of the columns
2 colnames(iris.hex)
3
4 names(iris.hex)
```

8.9 Getting Minimum and Maximum Values

Example in R

To view the maximum values for the real-valued columns in a data set, use `max()` with the name of the R reference object in the parentheses. To obtain the minimum values for the real-valued columns in a data set, use `min()` with the name of the R reference object in the parentheses.

```
1 min(prostate.hex$AGE)
2
3 max(prostate.hex$AGE)
```

8.10 Getting Quantiles

Example in R

To request quantiles for an H2O parsed data set, use `quantile()` with the name of the R reference object in the parentheses. To request a quantile for a single numerical column, use

`quantile(ReferenceObject$ColumnName)`, where

`ReferenceObject` represents the R reference object name and

`ColumnName` represents the name of the specified column. When you request for a full parsed data set consisting of a single column, `quantile()` displays a matrix with quantile information for the data set.

```
1 prosPath <- system.file("extdata", "prostate.csv",
2   package="h2o")
3 prostate.hex <- h2o.importFile(path = prosPath)
4 # Returns the percentiles at 0, 10, 20, ..., 100%
5 prostate.qs <- quantile(prostate.hex$PSA, probs =
6   (1:10)/10)
7 prostate.qs
8
9 # Take the outliers or the bottom and top 10% of data
10 PSA.outliers <- prostate.hex[prostate.hex$PSA <=
11   prostate.qs["10%"] | prostate.hex$PSA >=
12   prostate.qs["90%"],]
13
14 # Check that the number of rows return is about 20% of
15 the original data
16 nrow(prostate.hex)
17
18 nrow(PSA.outliers)
19
20 nrow(PSA.outliers)/nrow(prostate.hex)
```

8.11 Summarizing Data

Example in R

To generate a summary similar to the one in R for each of the columns in the data set, use `summary()` with the name of the R reference object in the parentheses. For continuous real functions, this produces a summary that includes information on quartiles, min, max, and mean. For factors, this produces information about counts of elements within each factor level.

```
1 summary(prostate.hex)
```

8.12 Summarizing Data in a Table

Example in R

To summarize data, use `h2o.table()`. Because H2O can handle larger data sets, it is possible to generate tables that are larger than R's capacity, so use caution when executing this command. To summarize multiple columns, use `head(h2o.table (ObjectName[, c(ColumnNumber,ColumnNumber)]))`, where `ObjectName` is the name of the object in R and `ColumnNumber` is the number of the column.

```
1 # Counts of the ages of all patients
2 head(as.data.frame(h2o.table(prostate.hex[, "AGE"])))
3
4
5 # Two-way table of ages (rows) and race (cols) of all
  patients
6 # Example: For the first row there is one count of a
  43 year old that's labeled as RACE = 0
7 h2o.table(prostate.hex[, c("AGE", "RACE")])
```

8.13 Generating Random Numbers

Example in R

To append a column of random numbers to an H2O data frame for testing/training data splits that are used for analysis and validation in H2O, use `h2o.runif()` with the name of the R reference object in the parentheses. This method is best for customized frame splitting; otherwise, use `h2o.splitFrame()`. However, `h2o.runif()` is not as fast or stable as `h2o.splitFrame()`.

```
1 prosPath <- system.file("extdata", "prostate.csv",
  package="h2o")
```

```

2 prostate.hex <- h2o.importFile(path = prosPath)
3
4 ## Creates object for uniform distribution on prostate
data set
5 s <- h2o.runif(prostate.hex)
6 summary(s) ## Summarize the results of h2o.runif
7
8 ## Create training set with threshold of 0.8
9 prostate.train <- prostate.hex[s <= 0.8,]
10 ##Assign name to training set
11 prostate.train <- h2o.assign(prostate.train, "prostate
    .train")
12 ## Create test set with threshold to filter values
greater than 0.8
13 prostate.test <- prostate.hex[s > 0.8,]
14 ## Assign name to test set
15 prostate.test <- h2o.assign(prostate.test, "prostate.
    test")
16 ## Combine results of test & training sets, then
display result
17 nrow(prostate.train) + nrow(prostate.test)
18
19 nrow(prostate.hex) ## Matches the full set

```

8.14 Splitting Frames

Example in R

To generate two subsets (according to specified ratios) from an existing H2O data set for testing/training, use `h2o.splitFrame()`, which returns contiguous sections of the data without random sampling.

```

1 # Splits data in prostate data frame with a ratio of
0.75
2 prostate.split <- h2o.splitFrame(data = prostate.hex ,
    ratios = 0.75)
3 # Creates training set from 1st data set in split
4 prostate.train <- prostate.split[[1]]
5 # Creates testing set from 2nd data set in split
6 prostate.test <- prostate.split[[2]]

```

8.15 Getting Frames

Example in R

To create a reference object to the data frame in H2O, use `h2o.getFrame()`. This is helpful for switching between the web UI and the R API or for multiple users accessing the same H2O instance. The following example assumes `prostate.hex` is in the key-value (KV) store.

```
1 prostate.hex <- h2o.getFrame(id = "prostate.hex")
```

8.16 Getting Models

Example in R

To create a reference object for the model in H2O, use `h2o.getModel()`. This is helpful for users that alternate between the web UI and the R API or multiple users accessing the same H2O instance. The following example assumes a GBM with the ID

`GBM_8e4591a9b413407b983d73fbd9eb44cf` is in the key-value (KV) store.

```
1 gbm.model <- h2o.getModel(model_id = "GBM_8  
e4591a9b413407b983d73fbd9eb44cf")
```


8.17 Listing H2O Objects

Example in R

To generate a list of all H2O objects generated during a session and each objects size in bytes, use `h2o.ls()`.

```
1 h2o.ls()
```

8.18 Removing H2O Objects

Example in R

To remove an H2O object on the server associated with the object in the R environment, use `h2o.rm()`. For optimal performance, we recommend removing the object from the R environment as well using `remove()`, with the name of the object in the parentheses. If you do not specify an R environment, then the current environment is used.

```
1 h2o.rm(ids = c("prostate.train", "prostate.test"))
2 h2o.ls()
```

8.19 Adding Functions

Example in R

User-defined functions no longer need to be added explicitly to the H2O instance. An R function can be defined and executed against an `H2OFrame`.

```
1 # Create an R functional expression
2 simpleFun <- function(x) { 2*x + 5 }
3 # Evaluate the expression across prostate's AGE column
4 calculated <- simpleFun(prostate.hex[, "AGE"])
5 h2o.cbind(prostate.hex[, "AGE"], calculated)
```

9 Running Models

This section describes how to run the following model types:

- Gradient Boosted Models (GBM)
- Generalized Linear Models (GLM)
- K-Means
- Principal Components Analysis (PCA)

as well as how to generate predictions.

9.1 Gradient Boosted Models (GBM)

Example in R

To generate gradient boosted models for developing forward-learning ensembles, use `h2o.gbm()`. In the parentheses, define `x` (the predictor variable vector), `y` (the integer or categorical response variable), the distribution type (multinomial is the default, gaussian is used for regression), and the name of the `H2OParsedData` object. For more information, use `help(h2o.gbm)`.

```
1 library(h2o)
2 h2o.init(nthreads = -1)
3 data(iris)
4 iris.hex <- as.h2o(iris,destination_frame = "iris.hex"
5   )
6 iris.gbm <- h2o.gbm(y = 1, x = 2:5, training_frame =
7   iris.hex, ntrees = 10,
8   max_depth = 3,min_rows = 2, learn_
9   rate = 0.2, distribution= "
10  gaussian")
11
12 # To obtain the Mean-squared Error by tree from the
13 # model object:
14 iris.gbm@model$scoring_history
```

Example in R

To generate a classification model that uses labels, use `distribution="multinomial"`:

```
1 iris.gbm2 <- h2o.gbm(y = 5, x = 1:4, training_frame =
2   iris.hex, ntrees = 15, max_depth = 5, min_rows =
3   2,
4   learn_rate = 0.01, distribution=
5   "multinomial")
6
7 iris.gbm2@model$training_metrics
```

9.2 Generalized Linear Models (GLM)

Example in R

Generalized linear models (GLM) are some of the most commonly-used models for many types of data analysis use cases. While some data can be analyzed

using general linear models, general linear models may not be as accurate if the variables are more complex. For example, if the dependent variable has a non-continuous distribution or if the effect of the predictors is not linear, generalized linear models will produce more accurate results than general linear models. Generalized Linear Models (GLM) estimate regression models for outcomes following exponential distributions in general. In addition to the Gaussian (i.e. normal) distribution, these include Poisson, binomial, gamma and Tweedie distributions. Each serves a different purpose, and depending on distribution and link function choice, it can be used either for prediction or classification. H2O's GLM algorithm fits the generalized linear model with elastic net penalties. The model fitting computation is distributed, extremely fast, and scales extremely well for models with a limited number (\sim low thousands) of predictors with non-zero coefficients. The algorithm can compute models for a single value of a penalty argument or the full regularization path, similar to `glmnet`. It can compute Gaussian (linear), logistic, Poisson, and gamma regression models. To generate a generalized linear model for developing linear models for exponential distributions, use `h2o.glm()`. You can apply regularization to the model by adjusting the `lambda` and `alpha` parameters. For more information, use `help(h2o.glm)`.

```

1 prostate.hex <- h2o.importFile(path = "https://raw.
  github.com/h2oai/h2o/master/smalldata/logreg/
  prostate.csv"
2                                     , destination_frame = "
                                     prostate.hex")
3
4 prostate.glm<-h2o.glm(y = "CAPSULE", x = c("AGE","RACE
  ","PSA","DCAPS"), training_frame = prostate.hex,
5                               family = "binomial", nfolds = 10, alpha
                               = 0.5)
6 prostate.glm@model$cross_validation_metrics

```

9.3 K-Means

Example in R

To generate a K-Means model for data characterization, use `h2o.kmeans()`. This algorithm does not rely on a dependent variable. For more information, use `help(h2o.kmeans)`.

```

1 h2o.kmeans(training_frame = iris.hex, k = 3, x = 1:4)

```

9.4 Principal Components Analysis (PCA)

Example in R

To map a set of variables onto a subspace using linear transformations, use `h2o.pcomp()`. This is the first step in Principal Components Regression. For more information, use `help(h2o.pcomp)`.

```
1 ausPath = system.file("extdata", "australia.csv",  
  package="h2o")  
2 australia.hex = h2o.importFile(path = ausPath)  
3  
4 australia.pca <- h2o.pcomp(training_frame = australia  
  .hex, transform = "STANDARDIZE", k = 3)  
5 australia.pca
```

9.5 Predictions

The following section describes some of the prediction methods available in H2O.

Predict: Generate outcomes of a data set with any model. Predict with GLM, GBM, Decision Trees or Deep Learning models.

Confusion Matrix: Visualize the performance of an algorithm in a table to understand how a model performs.

Area Under Curve (AUC): A graphical plot to visualize the performance of a model by its sensitivity, true positives and false positives to select the best model.

Hit Ratio: A classification matrix to visualize the ratio of the number of correctly classified and incorrectly classified cases.

PCA Score: Determine how well your feature selection fits a particular model.

Multi-Model Scoring: Compare and contrast multiple models on a data set to find the best performer to deploy into production.

Example in R

To apply an H2O model to a holdout set for predictions based on model results, use `h2o.predict()`. In the following example, H2O generates a model and then displays the predictions for that model. For classification, the predict column is the model's discrete prediction, based on maximum F1 by default; the individual class probabilities are the remaining columns in the data frame. It is common to utilize the p1 column for binary classification, if a raw probability is desired.

```
1 prostate.fit = h2o.predict(object = prostate.glm,  
    newdata = prostate.hex)  
2 prostate.fit
```

10 Appendix: Commands

The following section lists some common R commands by function and a brief description of each command.

10.1 Data Set Operations

Data Import/Export

`h2o.downloadCSV`: Download a H2O dataset to a CSV file on local disk.

`h2o.exportFile`: Export H2O Data Frame to a file.

`h2o.importFile`: Import a file from the local path and parse it.

`h2o.parseRaw`: Parse a raw data file.

`h2o.uploadFile`: Upload a file from the local drive and parse it.

Native R to H2O Coercion

`as.h2o`: Convert an R object to an H2O object.

H2O to Native R Coercion

`as.data.frame`: Check if an object is a data frame, or coerce it if possible.

Data Generation

`h2o.createFrame`: Create an H2O data frame, with optional randomization.

`h2o.runif`: Produce a vector of random uniform numbers.

`h2o.interaction`: Create interaction terms between categorical features of an H2O Frame.

Data Sampling/Splitting

`h2o.splitFrame`: Split an existing H2O data set according to user-specified ratios.

Missing Data Handling

`h2o.impute`: Impute a column of data using the mean, median, or mode.

`h2o.insertMissingValues`: Replaces a user-specified fraction of entries in a H2O dataset with missing values.

10.2 General Data Operations

Subscripting example to pull pieces from data object.

```

1  x[j]    ## note: chooses column J, not row J
2  x[i, j]
3  x[[i]]
4  x$name
5  x[i] <- value
6  x[i, j, ...] <- value
7  x[[i]] <- value
8  x$i <- value

```

Subsetting

`head`, `tail`: Return the First or Last Part of an Object

Concatenation

`c`: Combine Values into a Vector or List

`h2o.cbind`: Take a sequence of H2O datasets and combine them by column.

Data Attributes

`colnames`: Return column names for a parsed H2O data object.

`colnames<=`: Retrieve or set the row or column names of a matrix-like object.

`names`: Get the name of an object.

`names<=`: Set the name of an object.

`dim`: Retrieve the dimension of an object.

`length`: Get the length of vectors (including lists) and factors.

`nrow`: Return a count of the number of rows in an H2OParsedData object.

`ncol`: Return a count of the number of columns in an H2OParsedData object.

`h2o.anyFactor`: Check if an H2O parsed data object has any categorical data columns.

`is.factor`: Check if a given column contains categorical data.

Data Type Coercion

`as.factor`: Convert a column from numeric to factor.

`as.Date`: Converts a column from factor to date.

10.3 Methods from Group Generics

Math (H2O)

`abs`: Compute the absolute value of x .

`sign`: Return a vector with the signs of the corresponding elements of x (the sign of a real number is 1, 0, or -1 if the number is positive, zero, or negative, respectively).

`sqrt`: Computes the principal square root of x , \sqrt{x} .

`ceiling`: Take a single numeric argument x and return a numeric vector containing the smallest integers not less than the corresponding elements of x .

`floor`: Take a single numeric argument x and return a numeric vector containing the largest integers not greater than the corresponding elements of x .

`trunc`: Take a single numeric argument x and return a numeric vector containing the integers formed by truncating the values in x toward 0.

`log`: Compute logarithms (by default, natural logarithms).

`exp`: Compute the exponential function.

Math (generic)

`cummax`: Display a vector of the cumulative maxima of the elements of the argument.

`cummin`: Display a vector of the cumulative minima of the elements of the argument.

`cumprod`: Display a vector of the cumulative products of the elements of the argument.

`cumsum`: Display a vector of the cumulative sums of the elements of the argument.

`log10`: Compute common (i.e., base 10) logarithms

`log2`: Compute binary (i.e., base 2) logarithms.

`log1p`: Compute $\log(1+x)$ accurately also for $|x| \ll 1$.

`acos`: Compute the trigonometric arc-cosine.

`acosh`: Compute the hyperbolic arc-cosine.

`asin`: Compute the trigonometric arc-sine.

`asinh`: Compute the hyperbolic arc-sine.

`atan`: Compute the trigonometric arc-tangent.

`atanh`: Compute the hyperbolic arc-tangent.

`expm1`: Compute $\exp(x) - 1$ accurately also for $|x| \ll 1$.

`cos`: Compute the trigonometric cosine.

`cosh`: Compute the hyperbolic cosine.

`cospi`: Compute the trigonometric two-argument arc-cosine.

`sin`: Compute the trigonometric sine.

`sinh`: Compute the hyperbolic sine.

`sinpi`: Compute the trigonometric two-argument arc-sine.

`tan`: Compute the trigonometric tangent.

`tanh`: Compute the hyperbolic tangent.

`tanpi`: Compute the trigonometric two-argument arc-tangent.

`gamma`: Display the gamma function γx

`lgamma`: Display the natural logarithm of the absolute value of the gamma function.

`digamma`: Display the first derivative of the logarithm of the gamma function.

`trigamma`: Display the second derivative of the logarithm of the gamma function.

Math2 (H2O)

`round`: Round the values to the specified number of decimal places. The default is 0.

`signif`: Round the values to the specified number of significant digits.

Summary (H2O)

`max`: Display the maximum of all the input arguments.

`min`: Display the minimum of all the input arguments.

`range`: Display a vector containing the minimum and maximum of all the given arguments.

`sum`: Calculate the sum of all the values present in its arguments.

Summary (generic)

`prod`: Display the product of all values present in its arguments.

`any`: Given a set of logical vectors, determine if at least one of the values is true.

`all`: Given a set of logical vectors, determine if all of the values are true.

10.4 Other Aggregations

Non-Group Generic Summaries

`mean`: Generic function for the (trimmed) arithmetic mean.

`sd`: Calculate the standard deviation of a column of continuous real valued data.

`var`: Compute the variance of `x`.

`summary`: Produce result summaries of the results of various model fitting functions.

`quantile`: Obtain and display quantiles for H2O parsed data.

Row / Column Aggregation

`apply`: Apply a function over an H2O parsed data object (an array).

Group By Aggregation

`h2o.ddply`: Split H2O dataset, apply a function, and display results.

`h2o.group_by`: Apply an aggregate function to each group of an H2O dataset.

Tabulation

`h2o.table`: Use the cross-classifying factors to build a table of counts at each combination of factor levels.

10.5 Data Munging

General Column Manipulations

`is.na`: Display missing elements.

Element Index Selection

`h2o.which`: Display the row numbers for which the condition is true.

Conditional Element Value Selection

`h2o.ifelse`: Apply conditional statements to numeric vectors in H2O parsed data objects.

Numeric Column Manipulations

`h2o.cut`: Convert H2O Numeric Data to Factor.

Character Column Manipulations

`h2o.strsplit`: Splits the given factor column on the input split.

`h2o.tolower`: Change the elements of a character vector to lower case.

`h2o.toupper`: Change the elements of a character vector to lower case.

`h2o.trim`: Remove leading and trailing white space.

`h2o.gsub`: Match a pattern & replace all instances of the matched pattern with the replacement string globally.

`h2o.sub`: Match a pattern & replace the first instance of the matched pattern with the replacement string.

Factor Level Manipulations

`h2o.levels`: Display a list of the unique values found in a column of categorical data.

Date Manipulations

`h2o.month`: Convert the entries of a H2OParsedData object from milliseconds to months (on a 0 to 11 scale).

`h2o.year`: Convert the entries of a H2OParsedData object from milliseconds to years, indexed starting from 1900.

Matrix Operations

`%*%:` Multiply two matrices, if they are conformable. `t:` Given a matrix or `data.frame` `x`, `t` returns the transpose of `x`.

10.6 Data Modeling

Model Training: Supervised Learning

`h2o.deeplearning:` Perform Deep Learning neural networks on an `H2OParsedData` object.

`h2o.gbm:` Build gradient boosted classification trees and gradient boosted regression trees on a parsed data set.

`h2o.glm:` Fit a generalized linear model, specified by a response variable, a set of predictors, and a description of the error distribution.

`h2o.naiveBayes:` Build gradient boosted classification trees and gradient boosted regression trees on a parsed data set.

`h2o.prcomp:` Perform principal components analysis on the given data set.

`h2o.randomForest:` Perform random forest classification on a data set.

Model Training: Unsupervised Learning

`h2o.anomaly:` Detect anomalies in a H2O dataset using a H2O deep learning model with auto-encoding.

`h2o.deepfeatures:` Extract the non-linear features from a H2O dataset using a H2O deep learning model.

`h2o.kmeans:` Perform k-means clustering on a data set.

Grid Search

`h2o.grid:` Efficient method to build multiple models with different hyperparameters.

Model Scoring

`h2o.predict:` Obtain predictions from various fitted H2O model objects.

Classification Model Helpers

`h2o.accuracy:` Get the between cluster sum of squares.

`h2o.auc:` Retrieve the AUC (area under ROC curve).

`h2o.confusionMatrix`: Display prediction errors for classification data from a column of predicted responses and a column of actual (reference) responses in H2O.

`h2o.hit_ratio_table`: Retrieve the Hit Ratios. If `train`, `valid`, and `xval` parameters are `FALSE` (default), then the training Hit Ratios value is returned. If more than one parameter is set to `TRUE`, then a named list of Hit Ratio tables are returned, where the names are `train`, `valid`, or `xval`.

`h2o.performance`: Evaluate the predictive performance of a model via various measures.

Regression Model Helper

`h2o.mse`: Display the mean squared error calculated from a column of predicted responses and a column of actual (reference) responses in H2O.

Clustering Model Helper

`h2o.betweenss`: Get the between cluster sum of squares.

`h2o.centers`: Retrieve the Model Centers.

10.7 H2O Cluster Operations

H2O Key Value Store Access

`h2o.assign`: Assign H2O hex.keys to objects in their R environment.

`h2o.getFrame`: Get a reference to an existing H2O data set.

`h2o.getModel`: Get a reference to an existing H2O model.

`h2o.ls`: Display a list of object keys in the running instance of H2O.

`h2o.rm`: Remove H2O objects from the server where the instance of H2O is running, but does not remove it from the R environment.

H2O Object Serialization

`h2o.loadModel`: Load an `H2OModel` object from disk.

`h2o.saveModel`: Save an `H2OModel` object to disk to be loaded back into H2O using `h2o.loadModel`.

H2O Cluster Connection

`h2o.init (nthreads = -1)`: Connect to a running H2O instance using all CPUs on the host and check the local H2O R package is the correct version.

`h2o.shutdown`: Shut down the specified H2O instance. All data on the server will be lost!

H2O Load Balancing

`h2o.rebalance`: Rebalance (repartition) an existing H2O data set into given number of chunks (per Vec), for load-balancing across multiple threads or nodes.

H2O Cluster Information

`h2o.clusterInfo`: Display the name, version, uptime, total nodes, total memory, total cores and health of a cluster running H2O.

`h2o.clusterStatus`: Retrieve information on the status of the cluster running H2O.

H2O Logging

`h2o.clearLog`: Clear all H2O R command and error response logs from the local disk.

`h2o.downloadAllLogs`: Download all H2O log files to the local disk.

`h2o.logAndEcho`: Write a message to the H2O Java log file and echo it back.

`h2o.openLog`: Open existing logs of H2O R POST commands and error responses on the local disk.

`h2o.getLogPath`: Get the file path for the H2O R command and error response logs.

`h2o.startLogging`: Begin logging H2O R POST commands and error responses.

`h2o.stopLogging`: Stop logging H2O R POST commands and error responses.

H2O String Manipulation

`h2o.gsub`: String global substitution (all occurrences).

`h2o.strsplit`: String Split.

`h2o.sub`: String substitution (first occurrence).

`h2o.tolower`: Convert characters to lower case.

`h2o.toupper`: Convert characters to upper case.

`h2o.trim`: Trim spaces.

11 Authors

Spencer Aiello

Spencer comes from an unconventional background. After studying Physics and Math as an undergraduate at UCSC, he came to San Francisco to continue his education, earning his MS in Analytics from USF. Spencer has worked on a number of projects related to analytics in R, Python, Java, and SQL. At H2O, he works primarily on the R front-end and the backend Java R-interpreter.

Eric Eckstrand

Eric is a Quality and Performance Hacker at H2O.ai. Eric has formal education in computer science and systems engineering. Prior to joining H2O, Eric was a submariner in the US Navy. His roles included Reactor Controls Assistant and Communications Officer (USS Pittsburgh SSN-720) and Submarine Operations Officer (Destroyer Squadron One).

Anqi Fu

Anqi is a data hacker at H2O.ai. After completing her undergraduate at University of Maryland: College Park, she attended Stanford University. While there she earned her masters degree in Economics, and a second masters degree in Statistics. Her interests include machine learning and optimization.

Mark Landry

Mark Landry is a competition data scientist and product manager at H2O. He enjoys testing ideas in Kaggle competitions, where he is ranked in the top 100 in the world (top 0.03%) and well-trained in getting quick solutions to iterate over. Interests are multi-model architectures and helping the world make fewer models that perform worse than the mean.

Jessica Lanford

Jessica is a word hacker and seasoned technical communicator at H2O.ai. She brings our product to life by documenting the many features and functionality of H2O. Having worked for some of the top companies in technology including Dell, AT&T, and Lam Research, she is an expert at translating complex ideas to digestible articles.

Patrick Aboyoun

Patrick is a former math and data hacker at H2O.ai who has made a career out of creating and delivering software and training for data scientists, particularly those who love R. Patrick received an M.S. in Statistics from the University of Washington and a B.S. in Statistics from Carnegie Mellon University.

12 References

- H2O.ai Team. **H2O website**, 2015. URL <http://h2o.ai> H2O.ai Team.
- H2O documentation**, 2015. URL <http://docs.h2o.ai> H2O.ai Team.
- H2O GitHub Repository**, 2015. URL <https://github.com/h2oai>
- H2O.ai Team. **H2O Datasets**, 2015. URL <http://data.h2o.ai> H2O.ai Team.
- H2O JIRA**, 2015. URL <https://jira.h2o.ai> H2O.ai Team.
- H2O R Package Documentation**, 2015. URL http://h2o-release.s3.amazonaws.com/h2o/latest_stable_Rdoc.html H2O.ai Team.
- H2Ostream**, 2015. URL <https://groups.google.com/d/forum/h2ostream> Erin LeDell.
- R Ensemble Documentation**, 2014. URL <http://www.stat.berkeley.edu/~ledell/R/h2oEnsemble.pdf>
- H2O.ai Team. **h2o: R Interface for H2O**, 2015. URL <http://www.h2o.ai>. R package version 3.1.0.99999